

1. Problem Statement

Implement a generic tool using Apache Flink


- Matrix \rightarrow Tool \rightarrow Singular Value Decomposition

Not scoped on analyzing a specific dataset

2. Introduction

Goal of dimensionality reduction \rightarrow compress data

- syntactically (bytes) & semantically (concepts)



Through Singular Value Decomposition (SVD)

- $A = U \Sigma V^T$
- U, V orthonormal eigenvectors of AA^T, A^TA
- Σ eigenvalues of U, V in descending order

$$\begin{bmatrix} \text{Users} & \text{Movies} \end{bmatrix} = \begin{bmatrix} \text{Users} & \text{Genres} \end{bmatrix} \times \begin{bmatrix} \text{Genres} & \text{Movies} \end{bmatrix} \times \begin{bmatrix} \text{Genres} & \text{Genres} \end{bmatrix}$$

Reduction: only keep k of Σ

Numerical Calculation

- $\text{Lanczos}(A) = (U, \text{TriDiag})$
 - TriDiag is a symmetric, tridiagonal auxiliary matrix
- $\text{EigenDecomposition}(\text{TriDiag}) = (\Sigma, V)$

Since TriDiag is small, only parallelize Lanczos

3. Lanczos

Iterative Algorithm – *not* embarrassingly parallel

In each Iteration from 1 to k produce next

- $u_i = A \times u_{i-1} \rightarrow U$ (plus orthonormalization)
- u_0 is random / uniformly chosen
- $(a_i, b_i) \rightarrow \text{TriDiag}$ (a diagonal, b off diagonal)

The only parallelizable part is $A \times u_{i-1}$

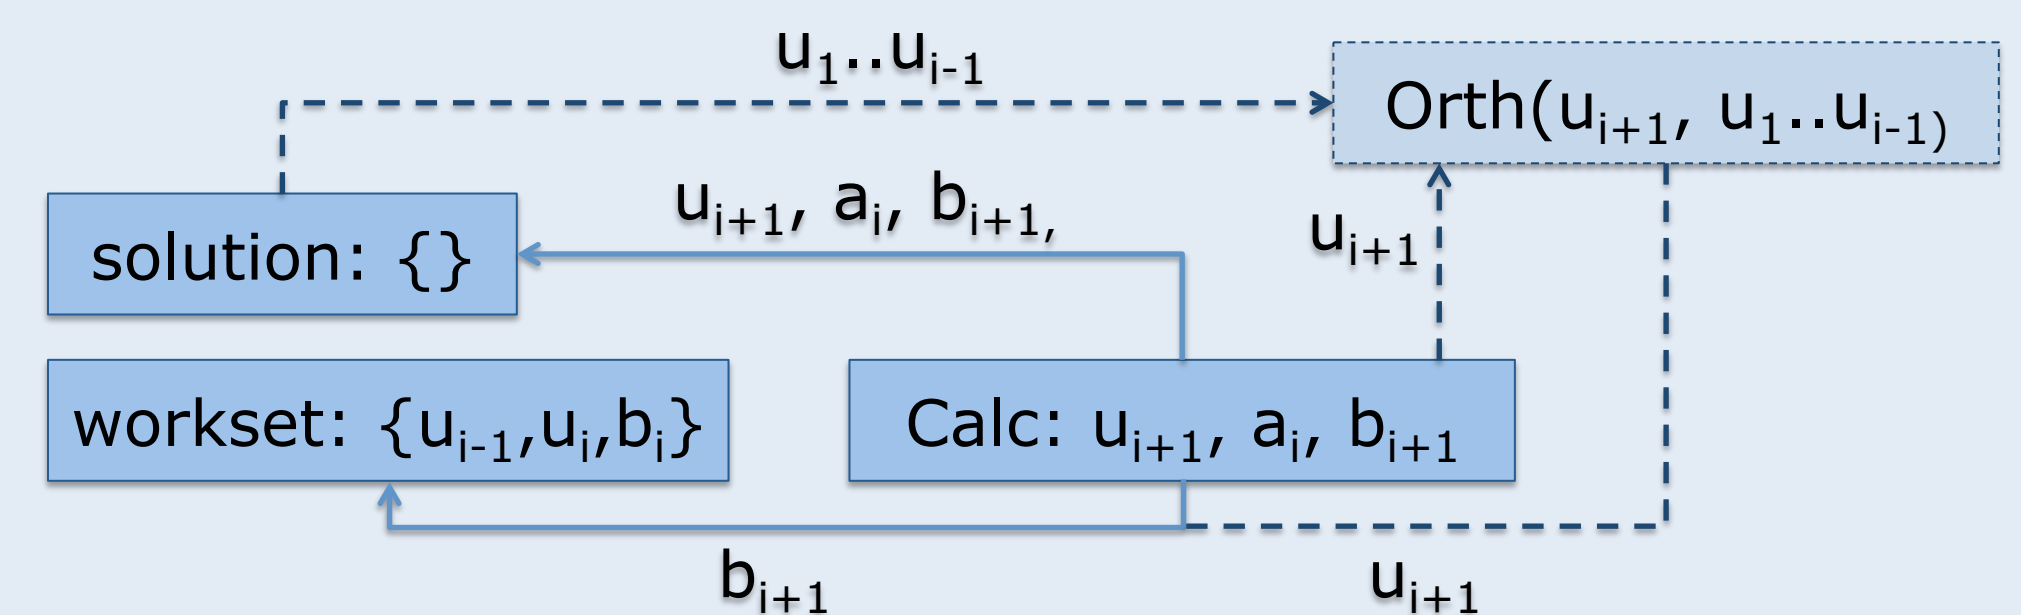
4. Approaches

- Delta Iterations
 - data flow from scratch optimizable by Flink
- Exploiting Mahout's Hadoop-Lanczos solver
 - Generic implementation (using interfaces)
 - By implementing interfaces with Flink
 - Each step one Hadoop-Job
- Iterative dataflow construction
 - Less optimizable by Flink

5. Implementations

1. Delta Iterations

- Requires to hold Basis + TriDiag in one DataSet
- $\text{Tuple4}(\text{id}, \text{row}, \text{col}, \text{val}) \rightarrow \text{„LanczosPlasma“}$



„Nested iterations are currently not supported“

2. Exploiting Mahout's Lanczos solver

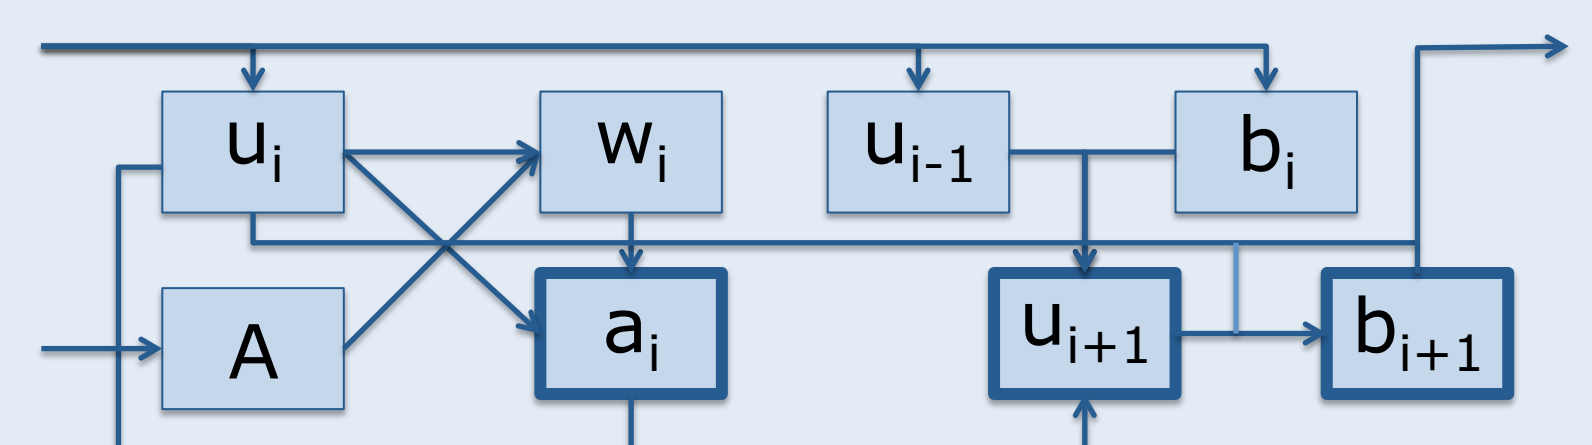
- relying on flat data types:

```
double alpha = currentVector.dot(nextVector);
```

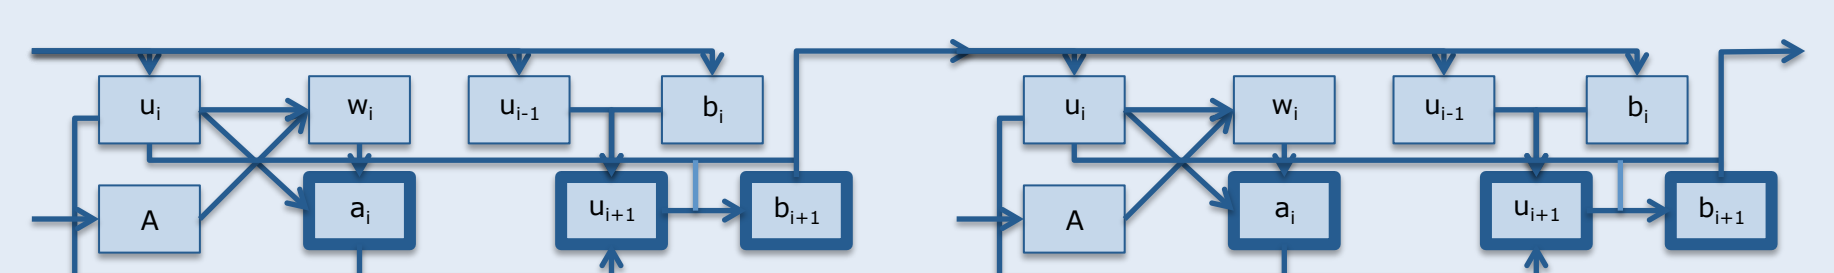
- Would require materialization in each iteration

3. Iterative dataflow construction

- One iteration step modelled as data flow



- Iteration is modelled by concatenating steps



- Compiler/Optimizer overloaded at 6 iterations

6. Results

- Idiomatic solution (delta iteration) not possible with Apache Flink's programming model (yet)
- Iterative data flow construction practically infeasible due to resulting DAG complexity
- Building on Mahout initially discarded due to intermediate materialization
 - \rightarrow only remaining approach that could work

7. Conclusion

- No satisfying solution possible with Flink 0.8.0
- Workaround by intermediate materialization is not expected to perform significantly better than Mahout + Hadoop
- Shows limitations of Apache Flink for implementing iterative algorithms
- Wait until Flink supports nested iterations